

Multi-Armed Bandits

Efficient Algorithms for Learning with Semi-Bandit Feedback

Abdul, Adwait, Anirudh

30th May 2021

- 1 Motivation and Introduction
- 2 FPL with GR Algorithm
- 3 Regret analysis and Comparison with State of art for FPL with GR
- 4 CombLinTS Algorithm
- 5 Regret analysis and Comparison with State of art for CombLinTS
- 6 CombLinUCB Algorithm
- 7 Regret analysis and Comparison with State of art for CombLinUCB

Motivation and Introduction

Goals of this section

- Motivate using Online Least Cost Path problem.
- Understanding the Semi-Bandits setting.
- Understanding the Combinatorial problem.
- Understanding Linear Generalization.
- Formally introducing the problem.


Motivation

The Online Least Cost Path problem

- We have a (directed or undirected) graph G with
 - L = number of edge,
 - w_i = cost associate with the edge i .
 - S = start node and,
 - T = target node.
- The goal: Find the cheapest path from S to T w.r.to the cost associated with the path.

Motivation

The Online Least Cost Path problem - An example



- In the above graph $L = 5$.
- A path is denoted as a 5 – D vector.
 - Path SAT, $A_1 = \begin{bmatrix} 1 & 0 & 0 & 1 & 0 \end{bmatrix}$
 - Path SCT, $A_2 = \begin{bmatrix} 0 & 0 & 1 & 0 & 1 \end{bmatrix}$
- Set of valid paths, $\mathcal{A} = \{A_1, A_2\} \subseteq \{0, 1\}^L$

Motivation

The Online Least Cost Path problem - Cost computation

- Let us say we have decided to go with the path SAT.

1	0	0	1	0
---	---	---	---	---
- To compute the cost, we may get one of the following:

- 1 Full information setting:

$$w = \begin{array}{|c|c|c|c|c|} \hline 0.3 & 0.2 & 0.7 & 0.5 & 0.4 \\ \hline \end{array}$$

- 2 Semi-Bandits setting:

$$w = \begin{array}{|c|c|c|c|c|} \hline 0.3 & x & x & 0.5 & x \\ \hline \end{array}$$

- 3 Full-Bandits setting:

$$l = \begin{array}{|c|} \hline 0.8 \\ \hline \end{array}$$

- For the Full information and Semi-Bandits setting, the loss is

$$l = A \cdot w$$

Motivation

How is this different from the familiar MAB setting?

- In the MAB setting, we choose $i \in [L]$.
 - i can take one of L values.
- However, here we choose $A \in \mathcal{A} \subseteq \{0, 1\}^L$.
 - A can take one of potentially 2^L values.
- That is, we want to choose one or more arms in each round.
- In the Online Least Cost path problem,
Edge \equiv Arm
Path \equiv multiple Arms.
- This is an example of **Combinatorial Optimization** problem.

- An immediate issue with this setting is that applying algorithms such as $EXP - 3$ may not have a good regret bound, since both L and $|\mathcal{A}|$ can be very large in real life settings.
- Can we make it independent of L ?
 - Here's a sneak peak...

Motivation

Linear Generalization

- In this setting (and many other real life settings), we have access to **features** of edges, such as,
 - 1 Length,
 - 2 Traffic, and
 - 3 Road Quality.

	Length	Traffic	Road Quality
e_1	10	3	7
e_2	4	6	3
e_3	8	1	4
e_4	6	2	5
e_5	3	7	2

Motivation

Linear Generalization

- Hence, we know a (possibly imperfect) **generalization matrix** Φ .

In this case,

$$\Phi = \begin{bmatrix} 10 & 3 & 7 \\ 4 & 6 & 3 \\ 8 & 1 & 4 \\ 6 & 2 & 5 \\ 3 & 7 & 2 \end{bmatrix}$$

- We assume that the cost of an edge is a **linear combination** of its feature values.

$$\text{cost}(e) = \theta_1 \text{length}(e) + \theta_2 \text{traffic}(e) + \theta_3 \text{quality}(e) = \vec{\theta} \cdot \phi_e$$

- Now, the task is to estimate the entries in $\vec{\theta}$.
- Notice the advantage:

The values to be estimated equal the number of features, and not L .

Let's formalize what all we have seen till now.

Introduction

Formalism

A Combinatorial Optimization Problem

Can be represented as a triple (E, \mathcal{A}, w) where,

- E = Set of L arms.
- \mathcal{A} = Set of (some or all) subsets A of E with $|A| \leq K$.
- $w : E \rightarrow \mathbb{R}$ is the weight function.
- The total weight of $A \in \mathcal{A}$ is $\sum_{e \in A} w_e$.

Introduction

Formalism

Semi-Bandit Problem

The losses associated to only the chosen arms are seen.

Linear Generalization

For $e \in [L]$, $w(e) = \Phi_e \theta$,

where Φ_e is the feature vector of the arm e , and θ indicates the weights of the features.

Introduction

Formalism

The Combinatorial semi-bandits

It is an online learning problem where at each step the learning agent chooses a subset of arms A subject to combinatorial constraints, and then observes weights of the selected arms, $\{w(e) : e \in A\}$, and gets their sum $f(A, w)$ as a payoff.

Introduction

Formalism

The Goal

- The learner chooses a subset of arms A^t at round t . The loss suffered at this round $R_t = f(A^*, w_t) - f(A^t, w)$ (In the Reward setup).
- We want the cumulative Regret to be *good*,
And the algorithm to be **efficient** when applied in Large scale settings.

Efficient Algorithm for Learning with Semi-Bandit Feedback (Paper 1)

Learning with Semi-Bandit Feedback

- The paper considers the problem of online Combinatorial Optimization under semi-bandit feedback. (No linear Generalization)
- The paper assumes:
 - Finite Decision Set (potentially very large)
 - Efficient offline combinatorial optimization is possible
 - Elements of the decision set can be described with d -dimensional arrays with at most m non-zero entries.
 - For example a decision vector in which 3 arms out of possible 7 arms are played will look like $[1 \ 0 \ 0 \ 1 \ 1 \ 0 \ 0]$

General Protocol for online Combinatorial Optimization

Parameters: set of decision vectors $\mathcal{S} = \{v(1), v(2), \dots, v(N)\} \subseteq \{0, 1\}^d$, number of rounds T ;

For all $t = 1, 2, \dots, T$, **repeat**

1. The learner chooses a probability distribution p_t over $\{1, 2, \dots, N\}$.
2. The learner draws an action I_t randomly according to p_t . Consequently, the learner plays decision vector $V_t = v(I_t)$.
3. The environment chooses loss vector ℓ_t .
4. The learner suffers loss $V_t^\top \ell_t$.
5. The learner observes some feedback based on ℓ_t and V_t .

- Generic framework to accommodate a number of interesting problem instances such as path planning, ranking and matching problems, finding minimum weight spanning trees and cut sets etc.

Normal Bandit Setting

- Total N arms and we chose only one arm.
- A distribution p_t over the arms where $p_{t,i} = \mathcal{P}[I_t = i | \mathcal{F}_{t-1}]$ where \mathcal{F}_{t-1} is the history of the learners observation and choice made upto step $t - 1$.
- Most bandit algorithms rely on feeding some loss estimates to black box algorithm like Hedge etc. A common loss estimate used is
$$\hat{\ell}_{t,i} = \frac{\ell_{t,i}}{p_{t,i}} \mathbb{I}[I_t = i] \quad (\text{An unbiased estimate})$$
- Almost all existing algorithms use some form of the above loss estimate. But p_t is not readily available for all algorithms like FPL.
- Following loss estimate is proposed

Geometric Resampling

- In round t the learner draws $l_t \sim p_t$
 - For $n = 1, 2, \dots$
 - Let $n \leftarrow n + 1$
 - Draw $l'(n) \sim p_t$
 - If $l'(n) = l_t$, break
 - Let $K_t = n$
-
- $\hat{\ell}_{t,i} = \ell_{t,i} \mathbb{I}[l_t = i] K_t$ (Also unbiased)
 - The number samples can be unbounded, so we cap the number of samples by M and use $\hat{K}_t = \min(K_t, M)$
 - Introduces some bias, but if M is chosen appropriately the performance does not hurt much.

Generalizing Geometric Sampling to Semi-Bandit Feedback

- In round t again the learner draws $l_t \sim p_t$
- For each arm draw M samples (because anyway we are capping by M). We can effectively draw only M samples and use them for all the arms being played. So, draw M additional indices $l'_t(1), l'_t(2) \dots l'_t(M) \sim p_t$
- Define $K_{t,j} = \min\{1 \leq s \leq M : v_j(l'_t(s)) = v_j(l_t)\}$ for all arms.
- Use Loss estimate for each arm as $\hat{\ell}_{t,j} = K_{t,j} V_{t,j} \ell_{t,j}$
- Since $V_{t,j}$ is non-zero only for the arms being played the estimates are well defined.

●

Input: $\mathcal{S} = \{v(1), v(2), \dots, v(N)\} \subseteq \{0, 1\}^d$, $\eta \in \mathbb{R}^+$, $M \in \mathbb{Z}^+$;

Initialization: $\hat{L}(1) = \dots = \hat{L}(d) = 0$;

for $t=1,\dots,T$ do

Draw $Z(1), \dots, Z(d)$ independently from distribution $\text{Exp}(\eta)$;

Choose action $I = \arg \min_{i \in \{1, 2, \dots, N\}} \left\{ v(i)^\top (\hat{L} - Z) \right\};$

$$K(1) = \cdots = K(d) = M;$$
 $k = 0;$

```
/* Counter for reoccurred indices */
```

for $n=1, \dots, M-1$ do

```
/* Geometric Resampling */
```

Draw $Z'(1), \dots, Z'(d)$ independently from distribution $\text{Exp}(\eta)$;

$$I(n) = \arg \min_{i \in \{1, 2, \dots, N\}} \left\{ \mathbf{v}(i)^\top \left(\hat{\mathbf{L}} - \mathbf{Z}' \right) \right\};$$
for $j=1,\dots,d$ do

if $v(I(n))(j) = 1$ & $K(j) = M$ then

$$K(j) = n;$$
$$k = k + 1;$$

```
if  $k = \|v(I)\|_1$  then break;      /* All indices reoccurred */
```

end

end

end

for $j=1,\dots,d$ do $\hat{L}(j) = \hat{L}(j) + K(j)v(I)(j)\ell(j)$;

```
/* Update */
```

end

Computational Efficiency under Semi Bandit Feedback

- The expected number of times the algorithm draws an action up to time step T can be upper bounded by dT . This can be shown as
- For each co-ordinate that the original arm had 1, we take sampling until we get 1 in the same co-ordinate again. (Here we do not assume cutoff)
- Let $q_{t,k} = \mathbb{E}[V_{t,k} | \mathcal{F}_{t-1}]$
- At time step t for a given co-ordinate k with 1, the expected number of samples is $\frac{1}{q_{t,k}}$ while the probability of co-ordinate k being 1 is $q_{t,k}$
- So, expected number of samples is $\sum_{k=1}^d q_{t,k} \frac{1}{q_{t,k}} = d$. (For one round)
- For T rounds it will be dT .
- The expected running time is comforting.

Regret Bound

- The total expected regret of FPL with Geometric Resampling satisfies $R_n \leq \frac{m(\log d+1)}{\eta} + \eta m d T + \frac{dT}{eM}$ under semi-bandit information.
 - For the full information setting if $C_T = \sum_{t=1}^T \mathbb{E}[V_t^\top \ell_t]$ then the expected regret of FPL satisfies $R_n \leq \frac{m(\log d+1)}{\eta} + \eta m C_t$ under full information.
- For $\eta = \sqrt{\frac{\log d+1}{dT}}$ and $M \geq \frac{\sqrt{dT}}{eM(\sqrt{\log d+1})}$, $R_n = \mathcal{O}(m\sqrt{dT \log d})$
 - For $\eta = \sqrt{\frac{(\log d+1)}{mT}}$ $R_n = \mathcal{O}(m^{\frac{3}{2}} \sqrt{T(\log d+1)})$.

Comparison with State of Art

- Full information setting \rightarrow Optimal Regret $\rightarrow \mathcal{O}(m\sqrt{T \log d})$
FPL-with-GR achieves $\mathcal{O}(m^{\frac{3}{2}} \sqrt{T(\log d + 1)})$ off by a factor of \sqrt{m}
- Semi-bandit setting \rightarrow Optimal Regret $\mathcal{O}(\sqrt{mdT})$.
FPL-with-GR achieves $\mathcal{O}(m\sqrt{dT \log d})$ off by a factor of $\sqrt{m \log d}$

Conclusion and Future work

- The paper has introduced the first general efficient algorithm for online combinatorial optimization under semi-bandit feedback.
- It remains as an open problem whether the gaps we have shown can be closed for FPL-style algorithms.
- The most important open problem is the development efficient online linear optimization with full bandit feedback.

Efficient Learning in Large Scale Combinatorial Semi-Bandits (Paper 2)

Large Scale Combinatorial Semi-Bandits

- This paper looks at efficient algorithms for stochastic combinatorial semi-bandits with **linear generalization**
- We look at algorithms that have regret bound **INDEPENDENT OF L** (hence large scale)

Setup Used for Combinatorial optimization

Combinatorial setup

$$(E, \mathcal{A}, P)$$

- $E = \{1, \dots, L\} \equiv$ arms (the ground set)
- $\mathcal{A} \subseteq \{A \subseteq E : |A| \leq K\} \equiv$ allowed combinations of arms
- $P \equiv$ probability distribution over weights $w \in \mathbb{R}^L$ on E . ($\bar{w} = \mathbb{E}[w]$)
- After the arms A are pulled, we observe the individual return of each arm $\{w(e) : e \in A\}$
- $f(A, w) \equiv$ loss on pulling $A \in \mathcal{A}$

We assume linear generalization(possibly imperfect) and that the agent knows the generalization matrix $\Phi \in \mathbb{R}^{L \times d}$. If $\bar{w} \in \text{span}[\Phi]$ we call it as coherent learning case otherwise agnostic. WLOG $\text{rank}[\Phi] = d$.

- $R_t = f(A^*, w_t) - f(A^t, w)$ where $A^* = \text{ORACLE}(E, \mathcal{A}, \bar{w})$
- For fixed \bar{w} ,
 $R(T) = \sum_{t=1}^T \mathbb{E}[R_t | \bar{w}] \equiv \text{EXPECTED CUMULATIVE REGRET}$
where expectation is over random weights and possible randomization in the algorithm.
- for randomly generated \bar{w} ,
 $R_{\text{bayes}}(T) = \sum_{t=1}^T \mathbb{E}[R_t] \equiv \text{BAYES CUMULATIVE REGRET}$ where expectation is over random weights, possible randomization in the algorithm and also over \bar{w} .
- $\theta^* = \text{argmin}_{\theta} ||\bar{w} - \Phi\theta||$. Since $\text{rank}[\Phi] = d$, θ^* is uniquely defined. Also in coherent learning case $\bar{w} = \Phi\theta^*$.

ALGORITHMS

KALMAN FILTERING for updating parameters

Algorithm 1

Input: $\bar{\theta}_t, \Sigma_t, \sigma$, and feature-observation pairs $\{(\phi_e, \mathbf{w}_t(e)) : e \in A^t\}$

Initialize $\bar{\theta}_{t+1} \leftarrow \bar{\theta}_t$ and $\Sigma_{t+1} \leftarrow \Sigma_t$

for $k = 1, \dots, |A^t|$ **do**

 Update $\bar{\theta}_{t+1}$ and Σ_{t+1} as follows, where a_k^t is the k th element in A^t

$$\bar{\theta}_{t+1} \leftarrow \left[I - \frac{\Sigma_{t+1} \phi_{a_k^t} \phi_{a_k^t}^T}{\phi_{a_k^t}^T \Sigma_{t+1} \phi_{a_k^t} + \sigma^2} \right] \bar{\theta}_{t+1} + \left[\frac{\Sigma_{t+1} \phi_{a_k^t}}{\phi_{a_k^t}^T \Sigma_{t+1} \phi_{a_k^t} + \sigma^2} \right] \mathbf{w}_t(a_k^t) \quad \text{and} \quad \Sigma_{t+1} \leftarrow \Sigma_{t+1} - \frac{\Sigma_{t+1} \phi_{a_k^t} \phi_{a_k^t}^T \Sigma_{t+1}}{\phi_{a_k^t}^T \Sigma_{t+1} \phi_{a_k^t} + \sigma^2},$$

end for

Output: $\bar{\theta}_{t+1}$ and Σ_{t+1}

CombLinTS

Input: Combinatorial structure (E, \mathcal{A}) , generalization matrix $\Phi \in \mathbb{R}^{L \times d}$, algorithm parameters $\lambda, \sigma > 0$, oracle ORACLE

Initialize $\Sigma_1 \leftarrow \lambda^2 I \in \mathbb{R}^{d \times d}$ and $\bar{\theta}_1 = 0 \in \mathbb{R}^d$

for all $t = 1, 2, \dots, n$ **do**

 Sample $\theta_t \sim N(\bar{\theta}_t, \Sigma_t)$

 Compute $A^t \leftarrow \text{ORACLE}(E, \mathcal{A}, \Phi\theta_t)$

 Choose set A^t , and observe $w_t(e)$, $\forall e \in A^t$

 Compute $\bar{\theta}_{t+1}$ and Σ_{t+1} based on Algorithm 1

end for

CombLinTS: Key Points

- $\lambda \rightarrow$ inverse regularization parameter.
- smaller λ makes the covariance matrix \sum_t closer to 0.
- smaller $\lambda \Rightarrow$ narrower prior \Rightarrow insufficient exploration \Rightarrow degraded performance of CombLinTS.
- σ controls the decrease rate of covariance matrix Σ_t .
- Large σ will lead to slow learning and a smaller σ will make the algorithm quickly converge to some sub-optimal coefficient vector.

Regret Bound

- If $\bar{w} = \Phi\theta^*$, the prior on θ^* is $\mathcal{N}(0, \lambda^2 I)$, the noises $(w_t(e) - \bar{w}(e))$ are i.i.d sampled from $\mathcal{N}(0, \sigma^2)$ then CombLinTS guarantees:

$$R_{\text{Bayes}}(T) = \tilde{O}(K\lambda\sqrt{dT\min\{\ln(L), d\}})$$

- The conditions ensure it is a coherent gaussian case.
- The \tilde{O} notation hides the logarithmic factors.
- The regret bound is a minimum of two bounds. The first bound is L -dependent as $O(\sqrt{\ln(L)})$. The second bound is L -independent but is $\tilde{O}(d)$ instead of $\tilde{O}(\sqrt{d})$.

Combinatorial Linear UCB

CombLinUCB

Input: Combinatorial structure (E, \mathcal{A}) , generalization matrix $\Phi \in \mathbb{R}^{L \times d}$, algorithm parameters $\lambda, \sigma, c > 0$, oracle ORACLE

Initialize $\Sigma_1 \leftarrow \lambda^2 I \in \mathbb{R}^{d \times d}$ and $\bar{\theta}_1 = 0 \in \mathbb{R}^d$

for all $t = 1, 2, \dots, n$ **do**

 Define the UCB weight vector \hat{w}_t as

$$\hat{w}_t(e) = \langle \phi_e, \bar{\theta}_t \rangle + c \sqrt{\phi_e^T \Sigma_t \phi_e} \quad \forall e \in E$$

 Compute $A^t \leftarrow \text{ORACLE}(E, \mathcal{A}, \hat{w}_t)$

 Choose set A^t , and observe $w_t(e), \forall e \in A^t$

 Compute $\bar{\theta}_{t+1}$ and Σ_{t+1} based on Algorithm 1

end for

CombLinUCB: Key Points

- $\lambda \rightarrow$ inverse regularization parameter.
- σ controls the decrease rate of covariance matrix Σ_t .
- The constant c controls the degree of optimism.
- Small $c \Rightarrow$ algorithm might converge to some sub-optimal coefficient vector due to insufficient exploration.
- Large $c \Rightarrow$ excessive exploration and slow learning.

CombLinUCB: Regret Bound

Assuming P is a subset of $[0, 1]^L$, the stochastic item weights are statistically independent under P . Then for the coherent learning case i.e. $\bar{w} = \Phi\theta$ we have: For any λ , σ and $\delta \in (0, 1)$ and any c satisfying

$$c \geq \frac{1}{\sigma} \sqrt{dT \ln(1 + \frac{Tk\lambda^2}{d\sigma^2})} + 2 \ln(\frac{1}{\delta}) + \frac{\|\theta^*\|_2}{\lambda} \text{ we have:}$$

$$R(T) \leq 2cK\lambda \sqrt{\frac{dT \ln(1 + \frac{Tk\lambda^2}{d\sigma^2})}{\ln(1 + \frac{\lambda^2}{\sigma^2})}} + TK\delta$$

Specifically if we chose $\lambda = \sigma = 1$, $\delta = \frac{1}{nK}$ and c is the lower bound on its above condition then:

Regret Bound

$$R(T) = \tilde{O}(Kd\sqrt{T})$$

Comparison with state of art

Standard Result

Standard results have been set for the no generalisation case.

No generalization \Rightarrow lower bound of $\Omega(\sqrt{LKT})$

$\Rightarrow \Phi = \mathbf{I} \Rightarrow L = d \Rightarrow$ No generalization lower bound = $\Omega(\sqrt{KdT})$

CombLinTS

$\tilde{O}(\sqrt{K \min\{\ln(L), d\}})$ larger than the $\Omega(\sqrt{KdT})$

CombLinUCB

$\tilde{O}(\sqrt{Kd})$ larger than the $\Omega(\sqrt{KdT})$

Note

The $\Omega(\sqrt{d})$ and $\Omega(\sqrt{K})$ factors are due to linear generalization. Full tightness analysis has been left to future work.

Conclusion and Future Work

- This paper has introduced two learning algorithms CombLinTS and CombLinUCB for stochastic combinatorial semi bandits with linear generalization.
- The main contribution has been that the paper has successfully introduced L independent bounds which is highly useful for real life problems e.g Online advertisements (millions of users and products).
- This paper has left it open to derive bounds for the agnostic learning case.
- Another open problem is how to extend the results to combinatorial semi bandits with non-linear generalization.

TAKEAWAYS

What have we seen so far

What have we seen so far

- **Geometric Resampling:**
 - Don't need p_t explicitly;
 - Computationally efficient;
 - Fits into any semi bandit algorithm

What have we seen so far

- **Geometric Resampling:**
 - Don't need p_t explicitly;
 - Computationally efficient;
 - Fits into any semi bandit algorithm
- **Linear Generalization:**
 - Makes regret bound L independent

THANK YOU